

Practical work 09 – 14/11/2019

Conv Neural Networks with Keras - part 2

Objectives

The objective of this PW is to understand some more advanced methods to train and use Convolutional Neural Networks (CNN) including data augmentation and techniques of visualisation inside the network. Another objective is to review classical deep architectures and understand the recent strategies to compose such architectures.

Submission

- **Deadline** : Monday 25 November, 12am (noon)
- **Format** : Zip with report and/or iPython notebook.

Exercise 1 Data Augmentation

Use the notebook `CIFAR10CNN_from_raw_augmented_data_stud.ipynb` available on Moodle as starting point.

Data augmentation - online augmentation




FIGURE 1 – Online data augmentation pipeline

a) Train a CNN

Define a CNN with the following structure : `CONV(32F,same)-RELU-CONV(32F,same)-RELU-MAXP(2)-CONV(32F,same)-RELU-MAXP(2)-DENSE`. Train the network using 10 epochs and batches of 128 images. Use a `categorical_crossentropy` loss and the `adam` optimizer.

b) Train the CNN with data augmentation

Re-read the section of the slides explaining the principles of data augmentation for images. Keras allows you to use an *online* data augmentation strategy as illustrated on Figure 1. Using the example given for the FashionMNIST dataset (cf. slides), implement a similar data augmentation strategy for CIFAR10.

You may try with different strategies and hyperparameter values of the data augmentation tool of Keras.

- a) Report the accuracy on the train set and on the test set for your different experiments. Do you observe an improvement using data augmentation ?
- b) Compare the evolution of the loss through the training epochs, with and without using data augmentation. Comment your observations.
- c) If you tried with different data augmentation strategies, which one seems to give the best results ?

Exercise 2 Visualisation of Activations

The objective is here to visualise the different activation maps in the network previously trained. The Figure 2 illustrates the principle for the first CONV layer on the first 6 filters of a given network.

Using the best of your network previously trained on CIFAR10 in exercise 1, implement a visualizer for the activations at different layer outputs.

- a) Read again the example of the visualisation presented on slides 23-24 of the class.
- b) Implement a code to visualise all the filters at a given layer. Hints : use subplots to have a grid of images, use `for` loops to avoid code repetition.
- c) For a given input image (e.g. `X_train[12]`), visualise the different activations maps of your network : outputs of CONV, RELU, MAXP. Comment on what you see.

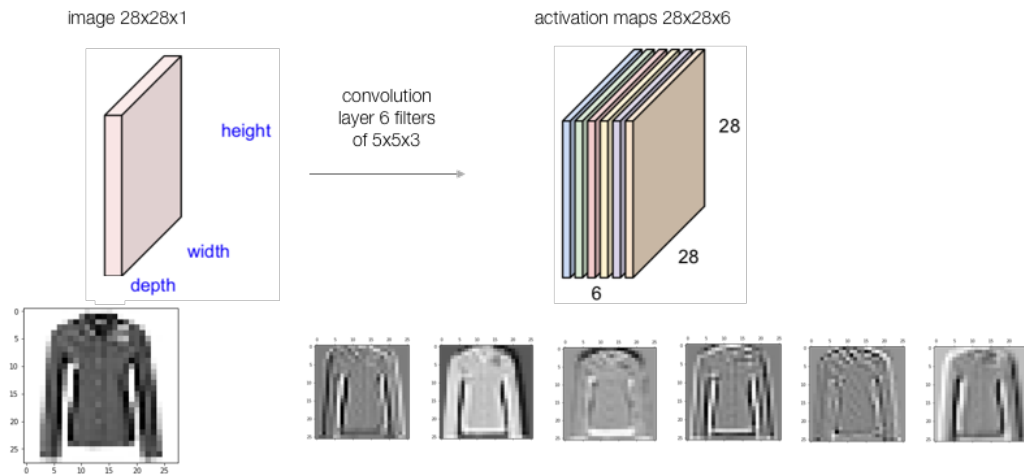


FIGURE 2 – Visualisation of activations of a CONV layer

Exercise 3 Optional : Deeper Models

Let's play here with a deeper CNN model on CIFAR10 using a structure inspired by VGGNet : `[[Conv2D → relu → BN]*3 → MaxPool2D → Dropout]*4 → Flatten → Dense → Dropout → Out`. In this structure, BN means Batch Normalisation. You can try different options but a configuration with blocks of 64 filters of size 3 with same padding and stride 1, max pooling of size 2 and stride 2, dropouts of 0.2 and a dense layer of 256 neurons should bring your performance around 80-85% on CIFAR10.

You probably now need to use a GPU to train and play with different settings. Report your best performances with and without data augmentation as in Exercise 1. What are your observations.





Exercise 4 Analysis of a Deep Architecture

In 2016, the **Inception-v4** architecture have been declared as outperforming ResNet and GoogleNet architecture on the ImageNet competition (see slide 61 of the class).

- Download the paper presenting the architecture at <https://arxiv.org/abs/1602.07261>.
- Read the paper up to the point you have a general understanding of their strategy¹.
- Re-explain in few phrases what you understood from the architecture doing comparison with the architectures presented in the class.

1. No need to understand all the details!

Exercise 5 Optional : Review Questions

- a) Explain the notion of hierarchical features with CNNs. 
- b) Explain 2 strategies to visualise the modelling taking place in CNNs. 
- c) What do we try to fight when using data augmentation? 
- d) What are the implementation strategies for data augmentation? 
- e) Explain the main differences for the deep architectures seen in class : AlexNet, VGGNet, GoogLeNet, ResNet. What were their intuitions when putting together such architectures?